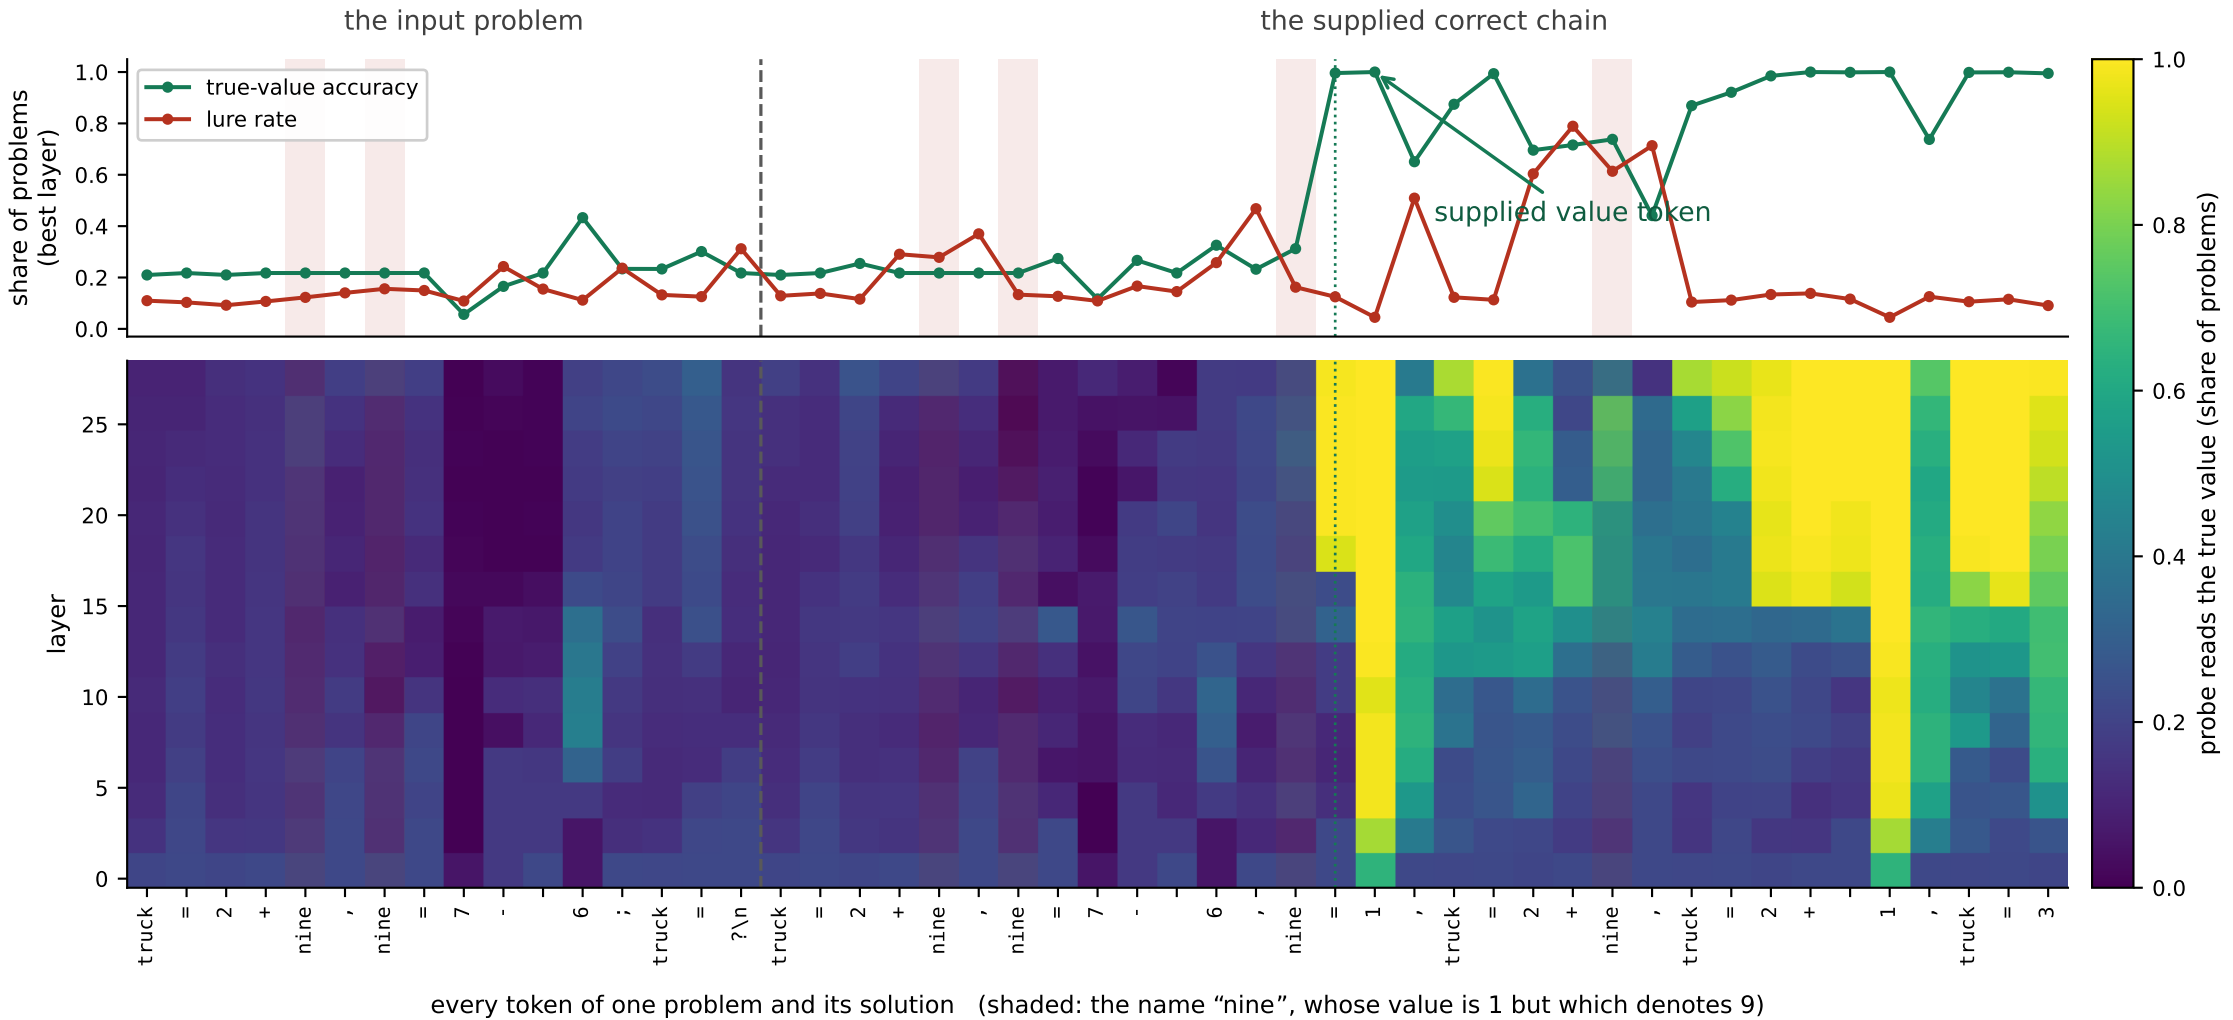

Llama-3.2-3B, with chain of thought: intermediate-value probe readouts
the input problem
the supplied correct chain
true-value accuracy
lure rate
share of problems (best layer)
supplied value token
layer
probe reads the true value (share of problems)
every token of one problem and its solution (shaded: the name “nine”, whose value is 1 but which denotes 9)