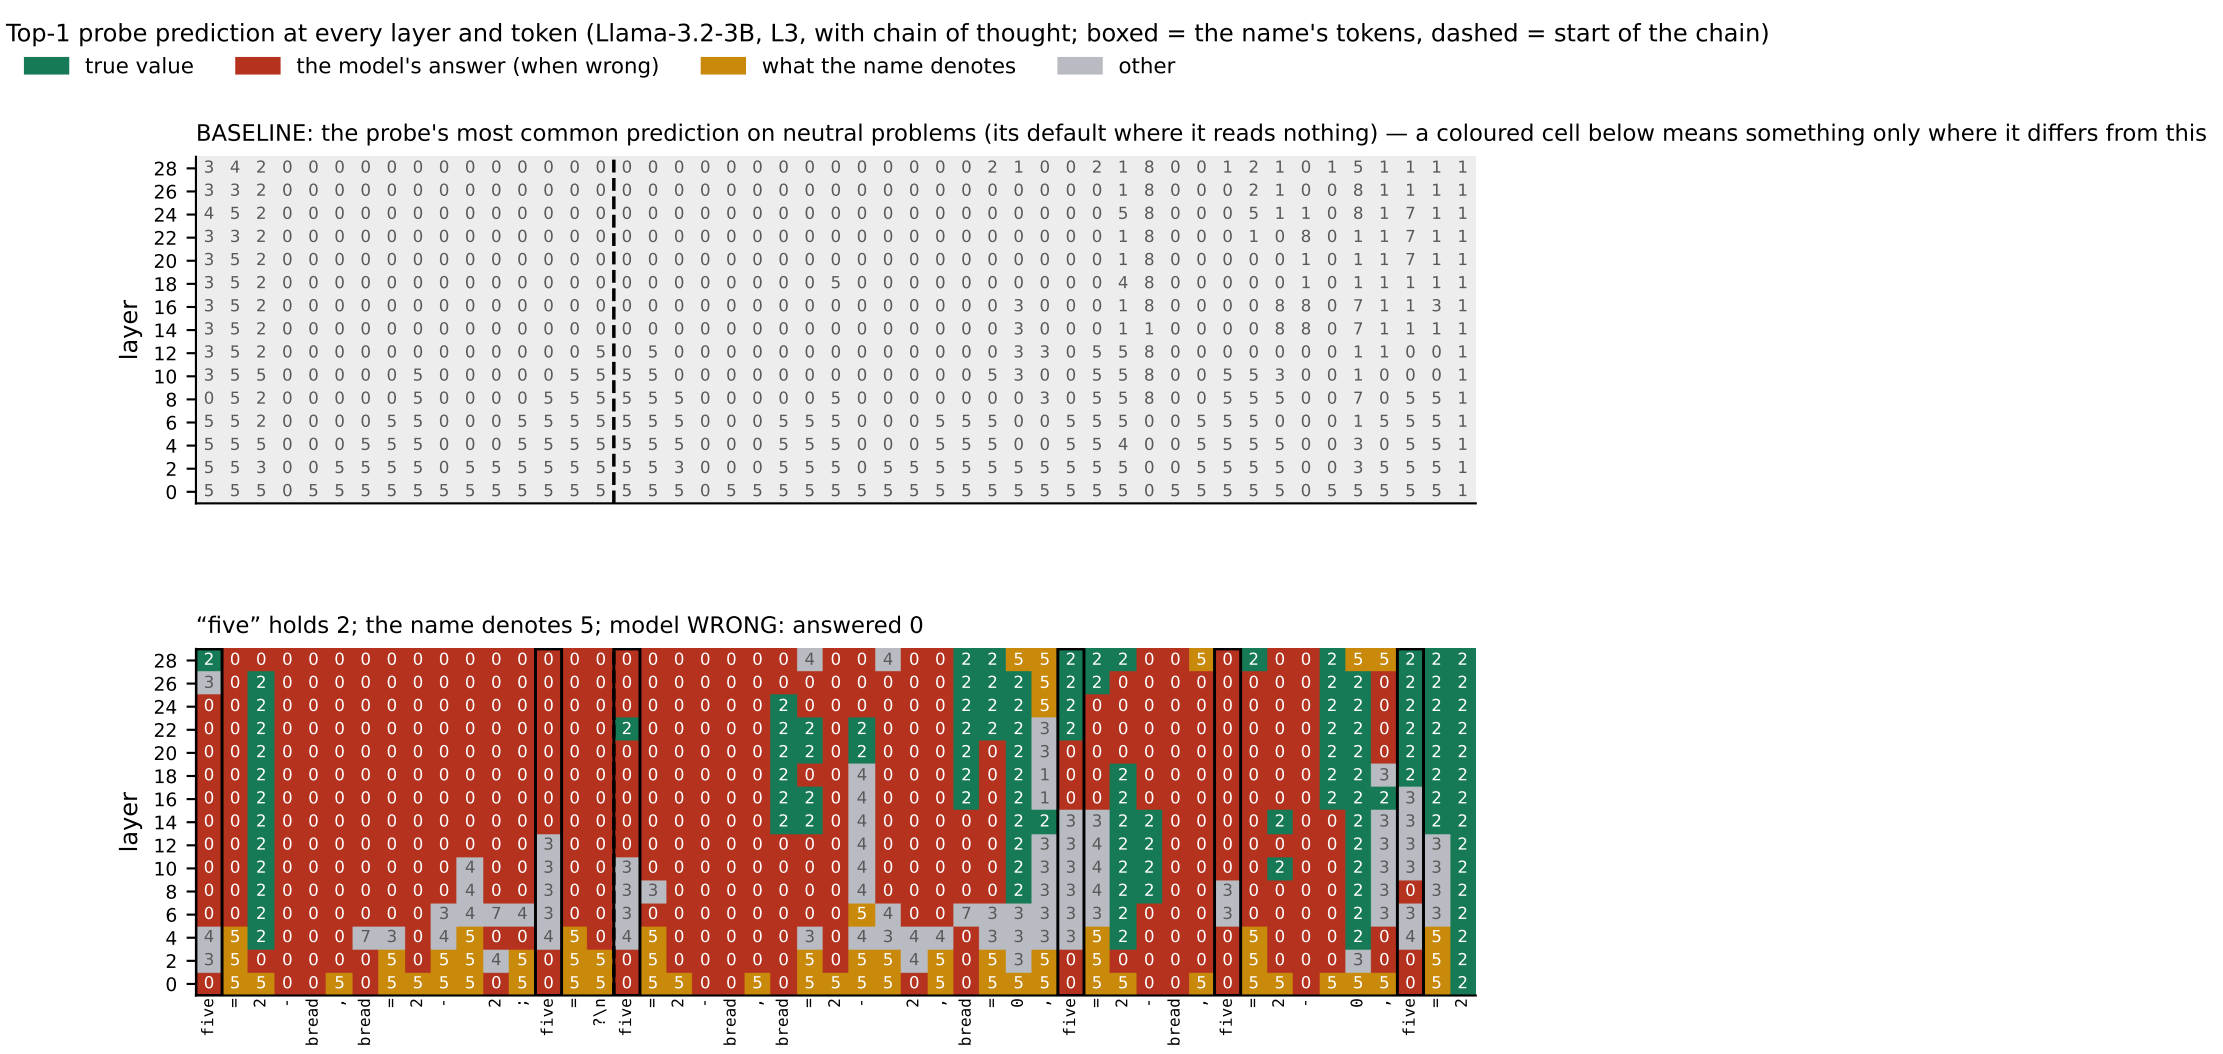

Top-1 probe prediction at every layer and token (Llama-3.2-3B, L3, with chain of thought; boxed = the name's tokens, dashed = start of the chain)
true value
the model's answer (when wrong)
what the name denotes
other
BASELINE: the probe's most common prediction on neutral problems (its default where it reads nothing) — a coloured cell below means something only where it differs from this
layer
“five” holds 2; the name denotes 5; model WRONG: answered 0
layer
five = 2 - bread , bread = 2 - 2 ; five = ?\n five = 2 - bread , bread = 2 - 2 , bread = 0 , five = 2 - bread , five = 2 - 0 , five = 2