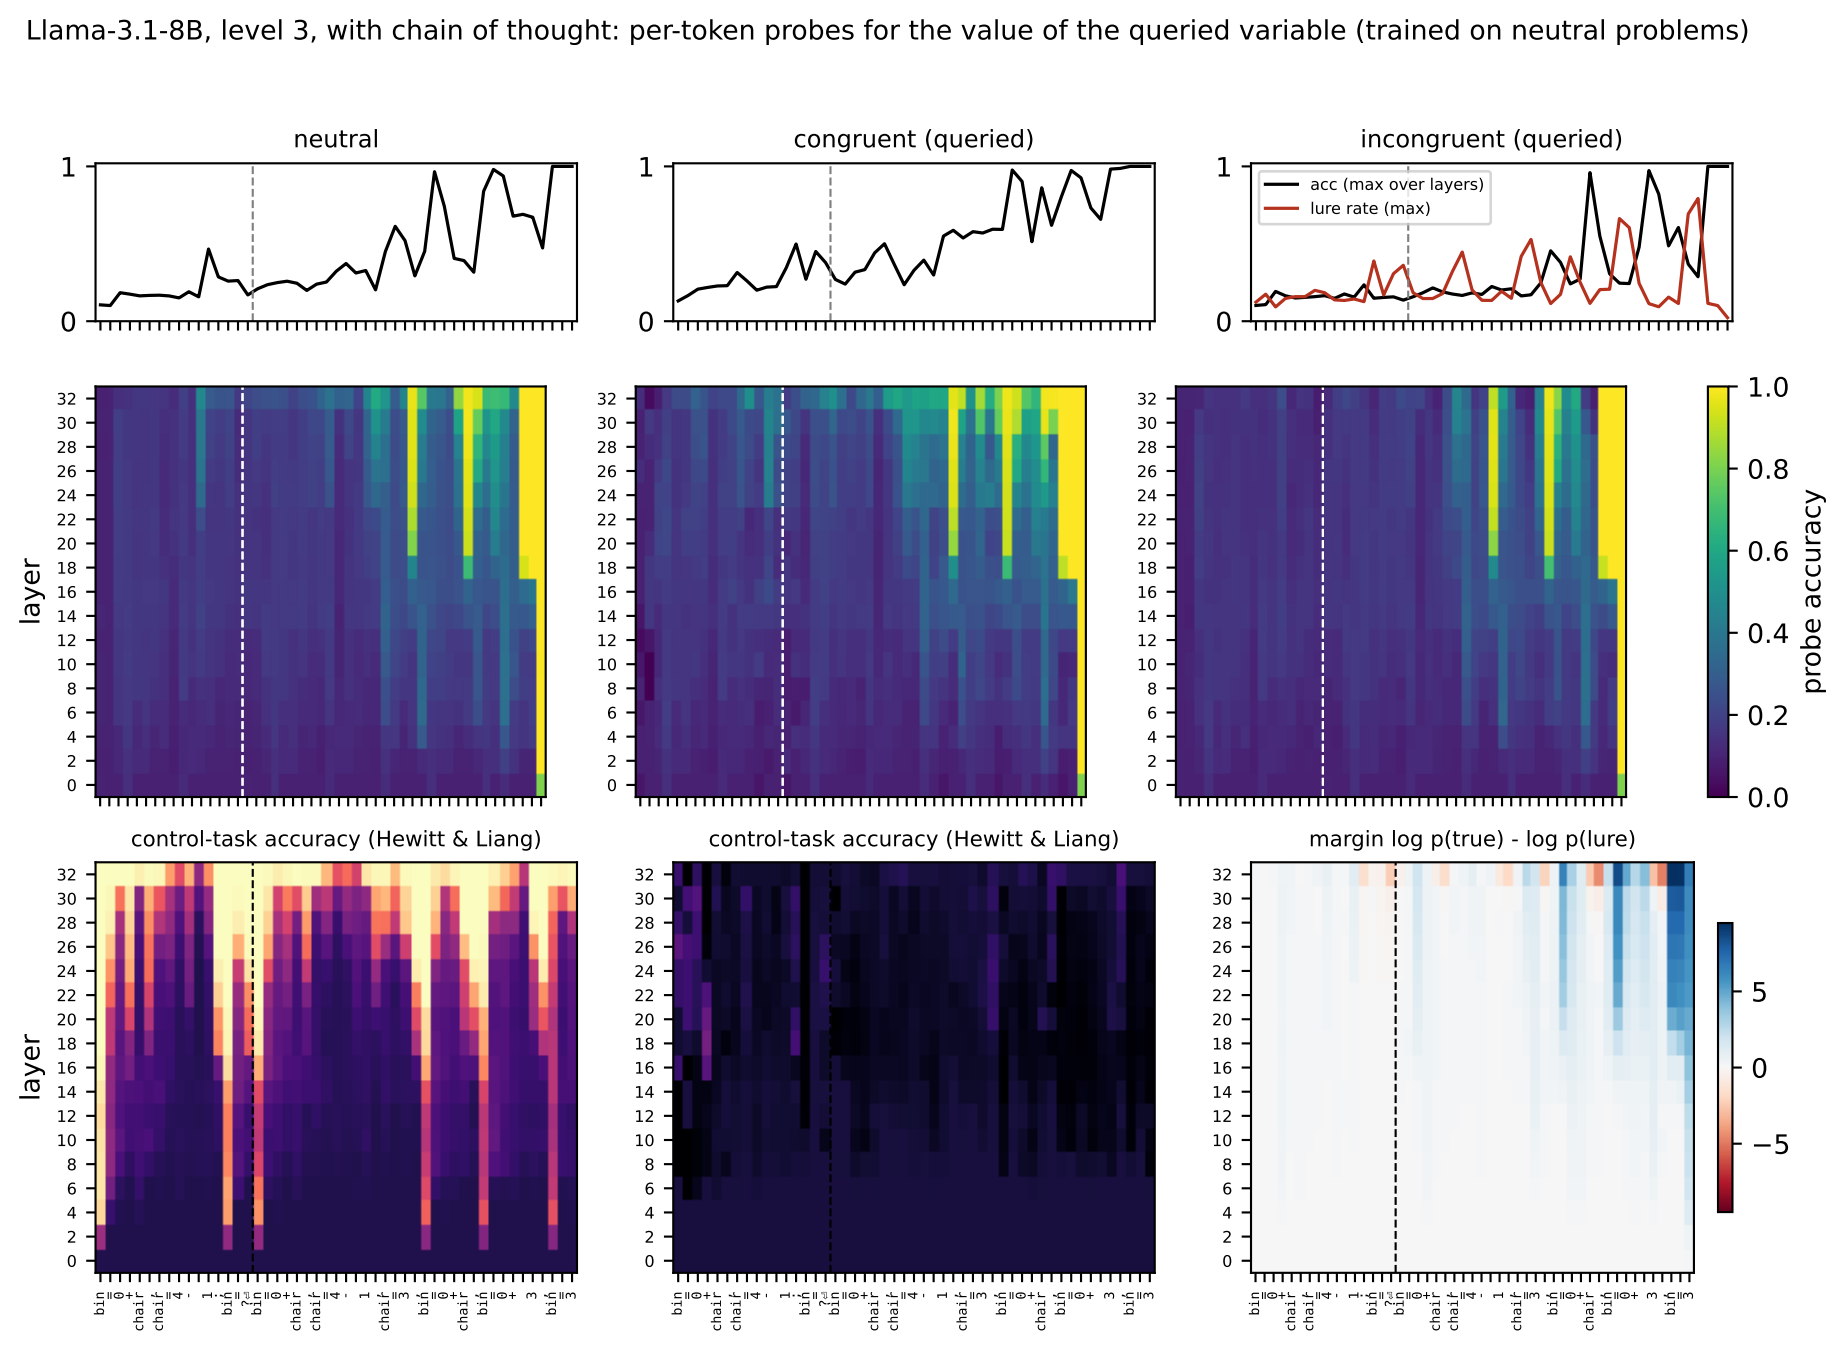

Llama-3.1-8B, level 3, with chain of thought: per-token probes for the value of the queried variable (trained on neutral problems)
neutral
congruent (queried)
incongruent (queried)
acc (max over layers)
lure rate (max)
layer
probe accuracy
control-task accuracy (Hewitt & Liang)
control-task accuracy (Hewitt & Liang)
margin log p(true) - log p(lure)