

(a) the manipulation

neutral

truck=2 + ant, ant=7 - 6; truck=?
the answer is 3

the name agrees

three=2 + ant, ant=7 - 6; three=?
the answer is 3, and the name says 3

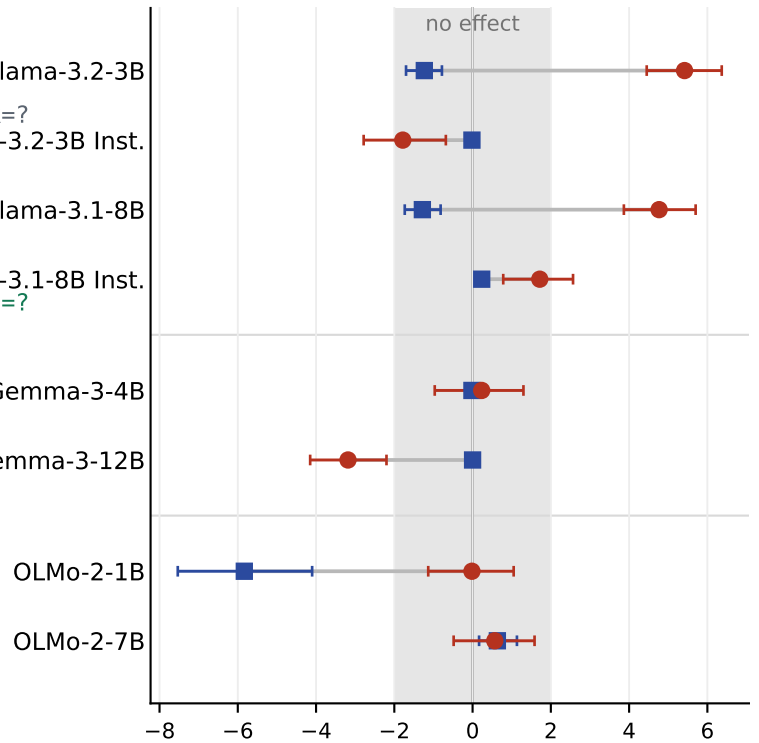
the name lies

four=2 + ant, ant=7 - 6; four=?
the answer is 3, but the name says 4

Same arithmetic every time.
Only the name of the variable
being asked about changes.

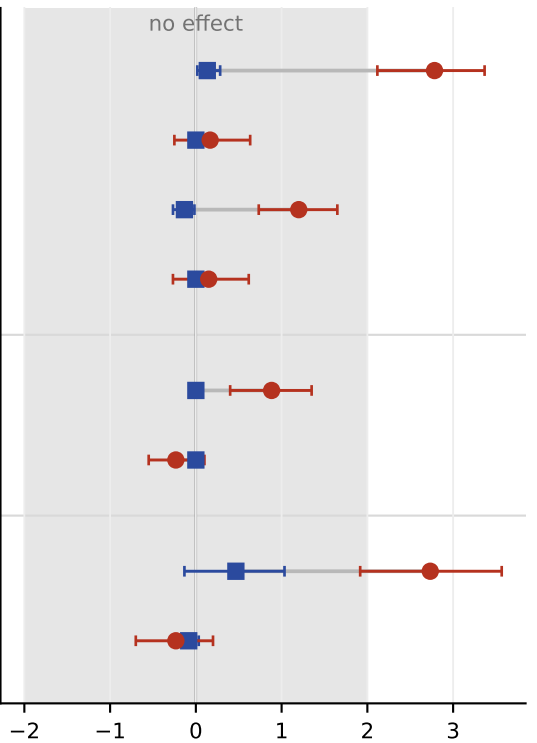
The model either writes out
the steps, ending
“four=2 + 1, four=3”,
or answers “four=3” directly.

(b) when the name says the right answer,
is the model more often right?



● no chain of thought ■ with chain of thought

when the name says a wrong number,
does the model answer that instead?



how much more often the answer is
the name's own number (points)

■ within 2 points of no effect

(c) what the model holds, token by token

