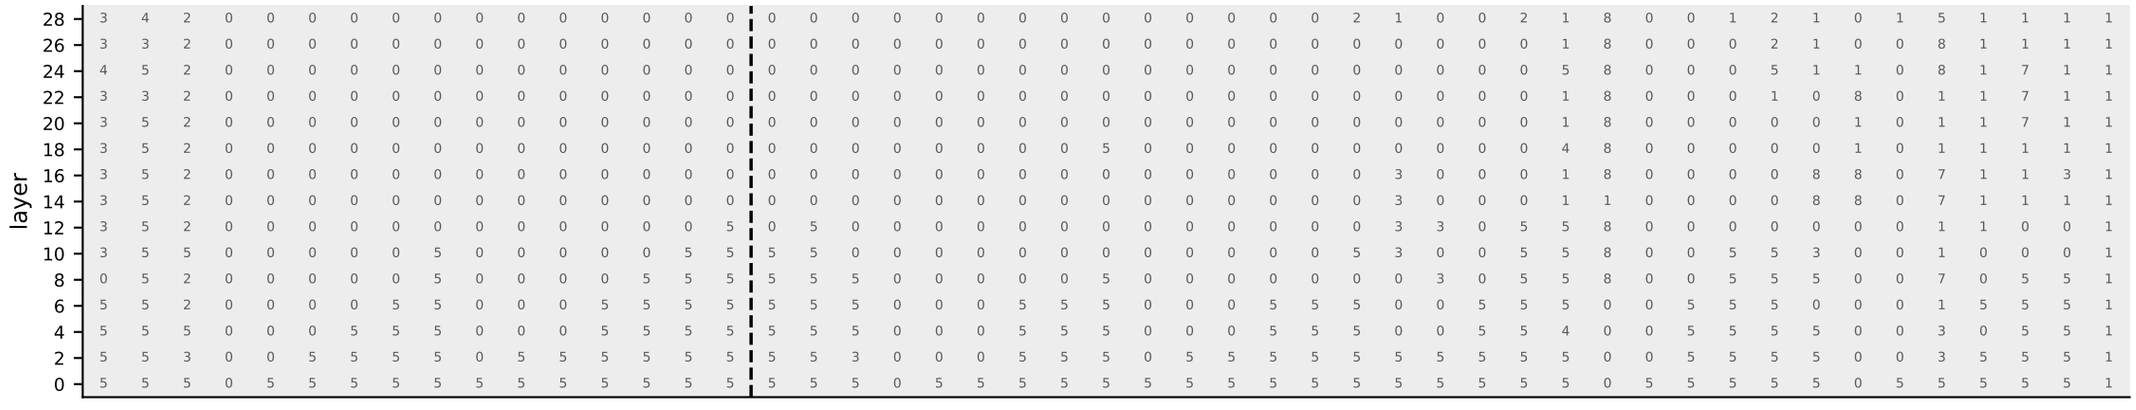
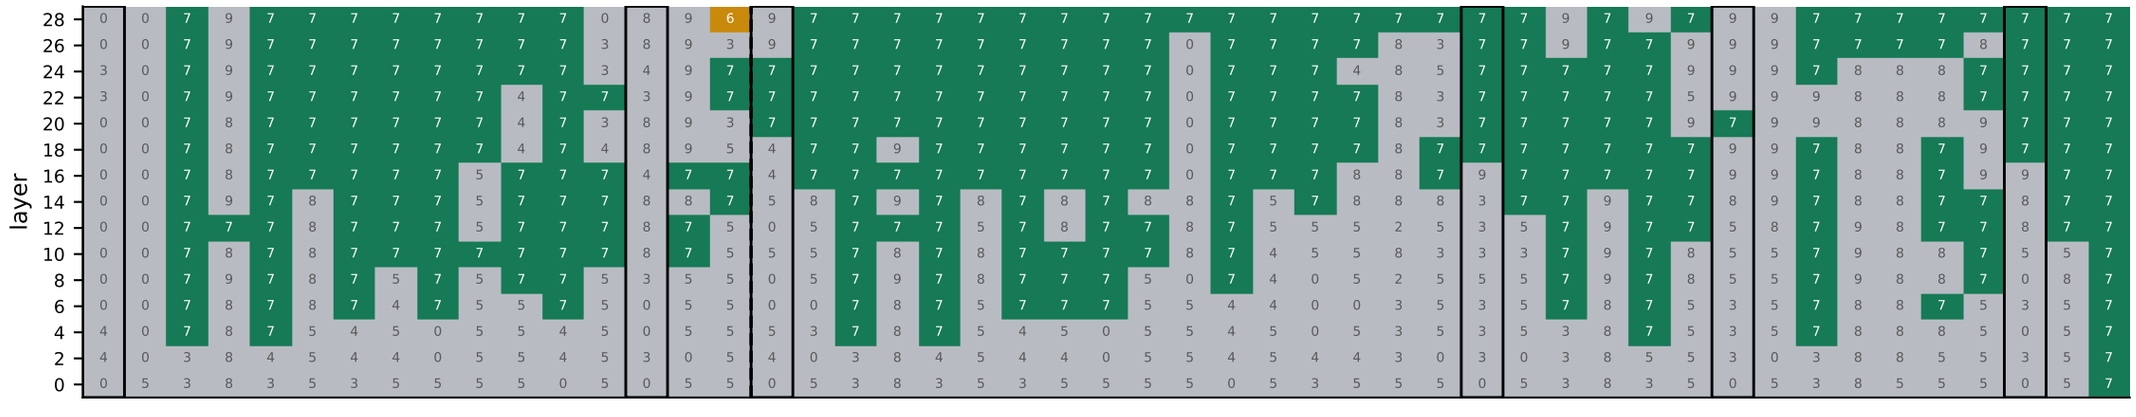
Top-1 probe prediction at every layer and token (Llama-3.2-3B, level 3, with chain of thought; boxed columns are the name's tokens, dashed line starts the chain)

the variable's true value · the value the model answered (when wrong) · the value the name denotes · any other digit

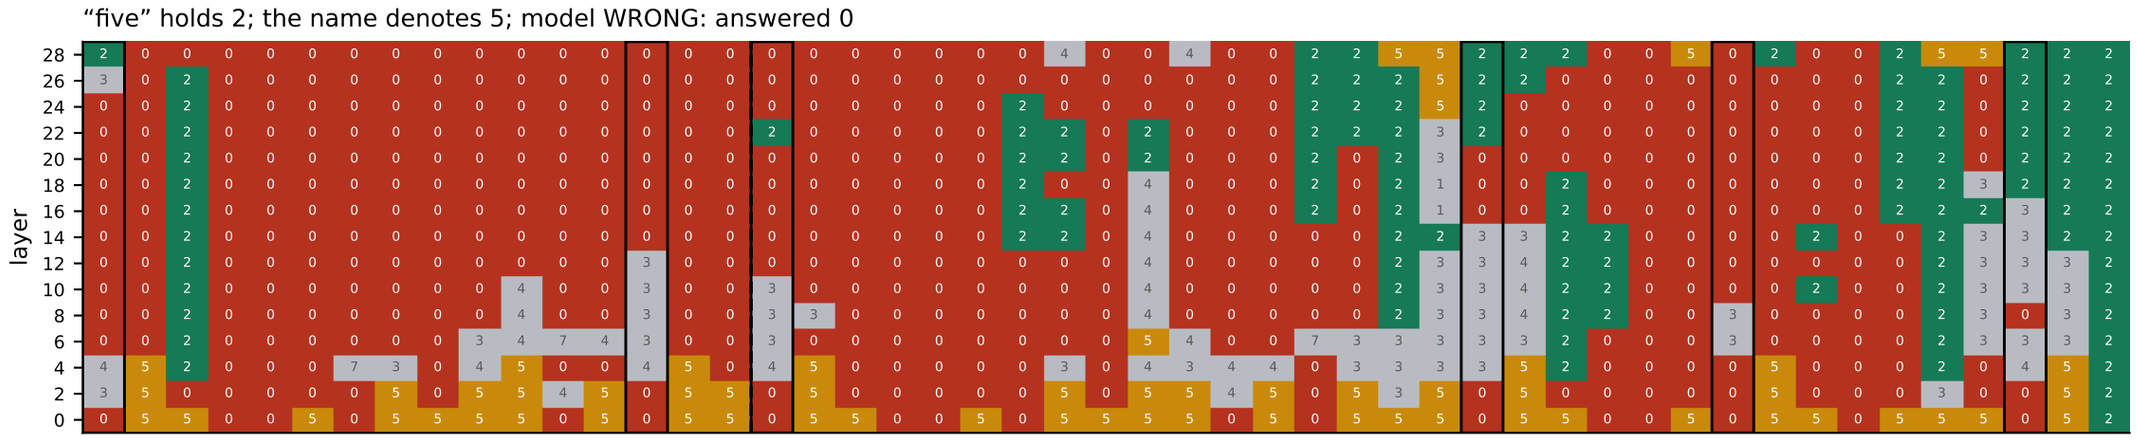
BASELINE: the probe's most common prediction on neutral problems (its default where it reads nothing) — a coloured cell below means something only where it differs from this




“six” holds 7; the name denotes 6; model WRONG: answered 11




“five” holds 2; the name denotes 5; model WRONG: answered 0




“nine” holds 6; the name denotes 9; model WRONG: answered 2

