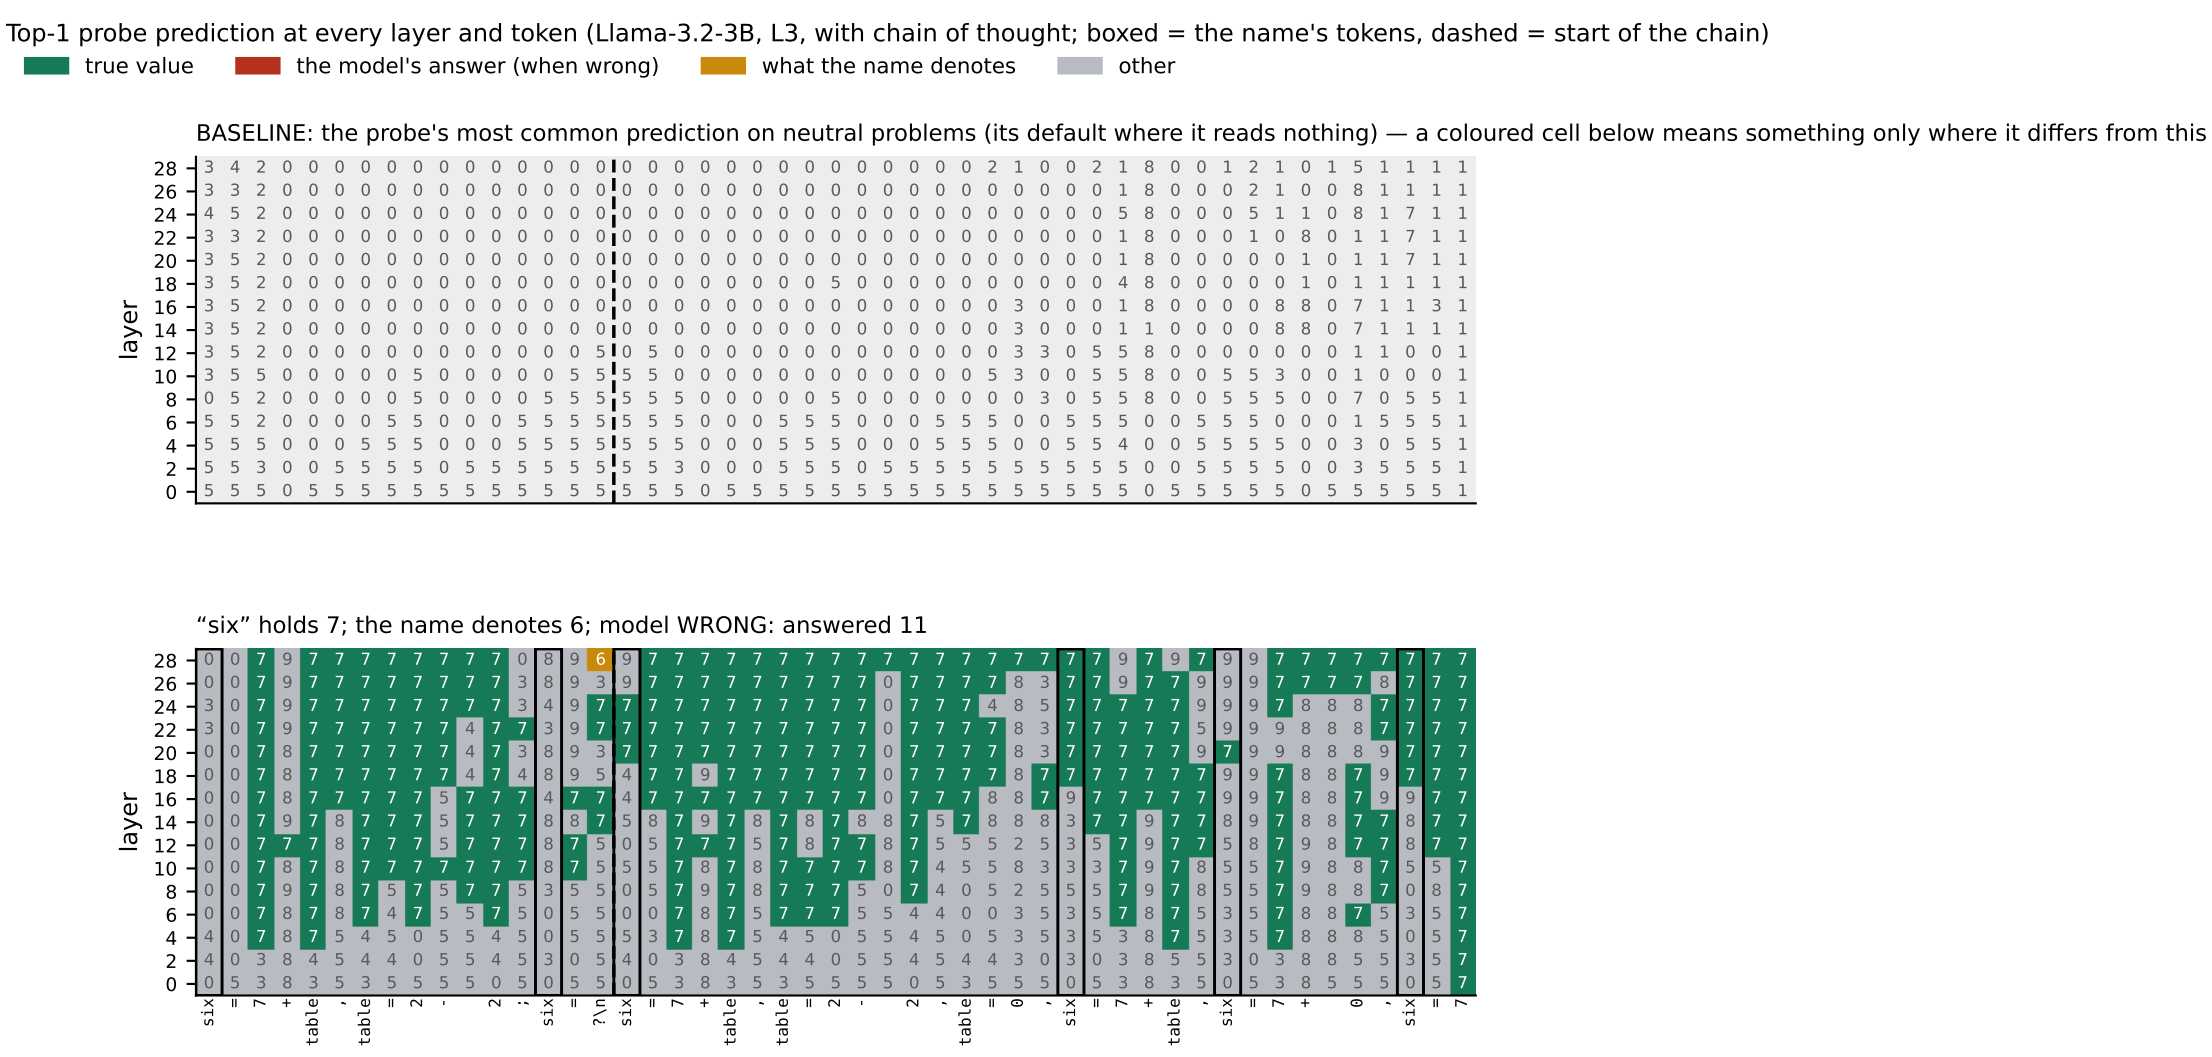

Top-1 probe prediction at every layer and token (Llama-3.2-3B, L3, with chain of thought; boxed = the name's tokens, dashed = start of the chain)
true value
the model's answer (when wrong)
what the name denotes
other
BASELINE: the probe's most common prediction on neutral problems (its default where it reads nothing) — a coloured cell below means something only where it differs from this
layer
“six” holds 7; the name denotes 6; model WRONG: answered 11
six = 7 + table , table = 2 - 2 ; six = ?\n six = 7 + table , table = 2 - 2 , table = 0 , six = 7 + table , six = 7 + 0 , six = 7