



Top-1 probe prediction at every layer and token (Llama-3.2-3B, level 3, with chain of thought; boxed columns are the name's tokens, dashed line starts the chain)

the variable's true value
 the value the model answered (when wrong)
 the value the name denotes
 any other digit


BASELINE: the probe's most common prediction on neutral problems (its default where it reads nothing) — a coloured cell below means something only where it differs from this



“four” holds 0; the name denotes 4; model WRONG: answered 4



"four" holds 3: the name denotes 4: model WRONG: answered 4



"four" holds 2: the name denotes 4: model WRONG: answered 4

