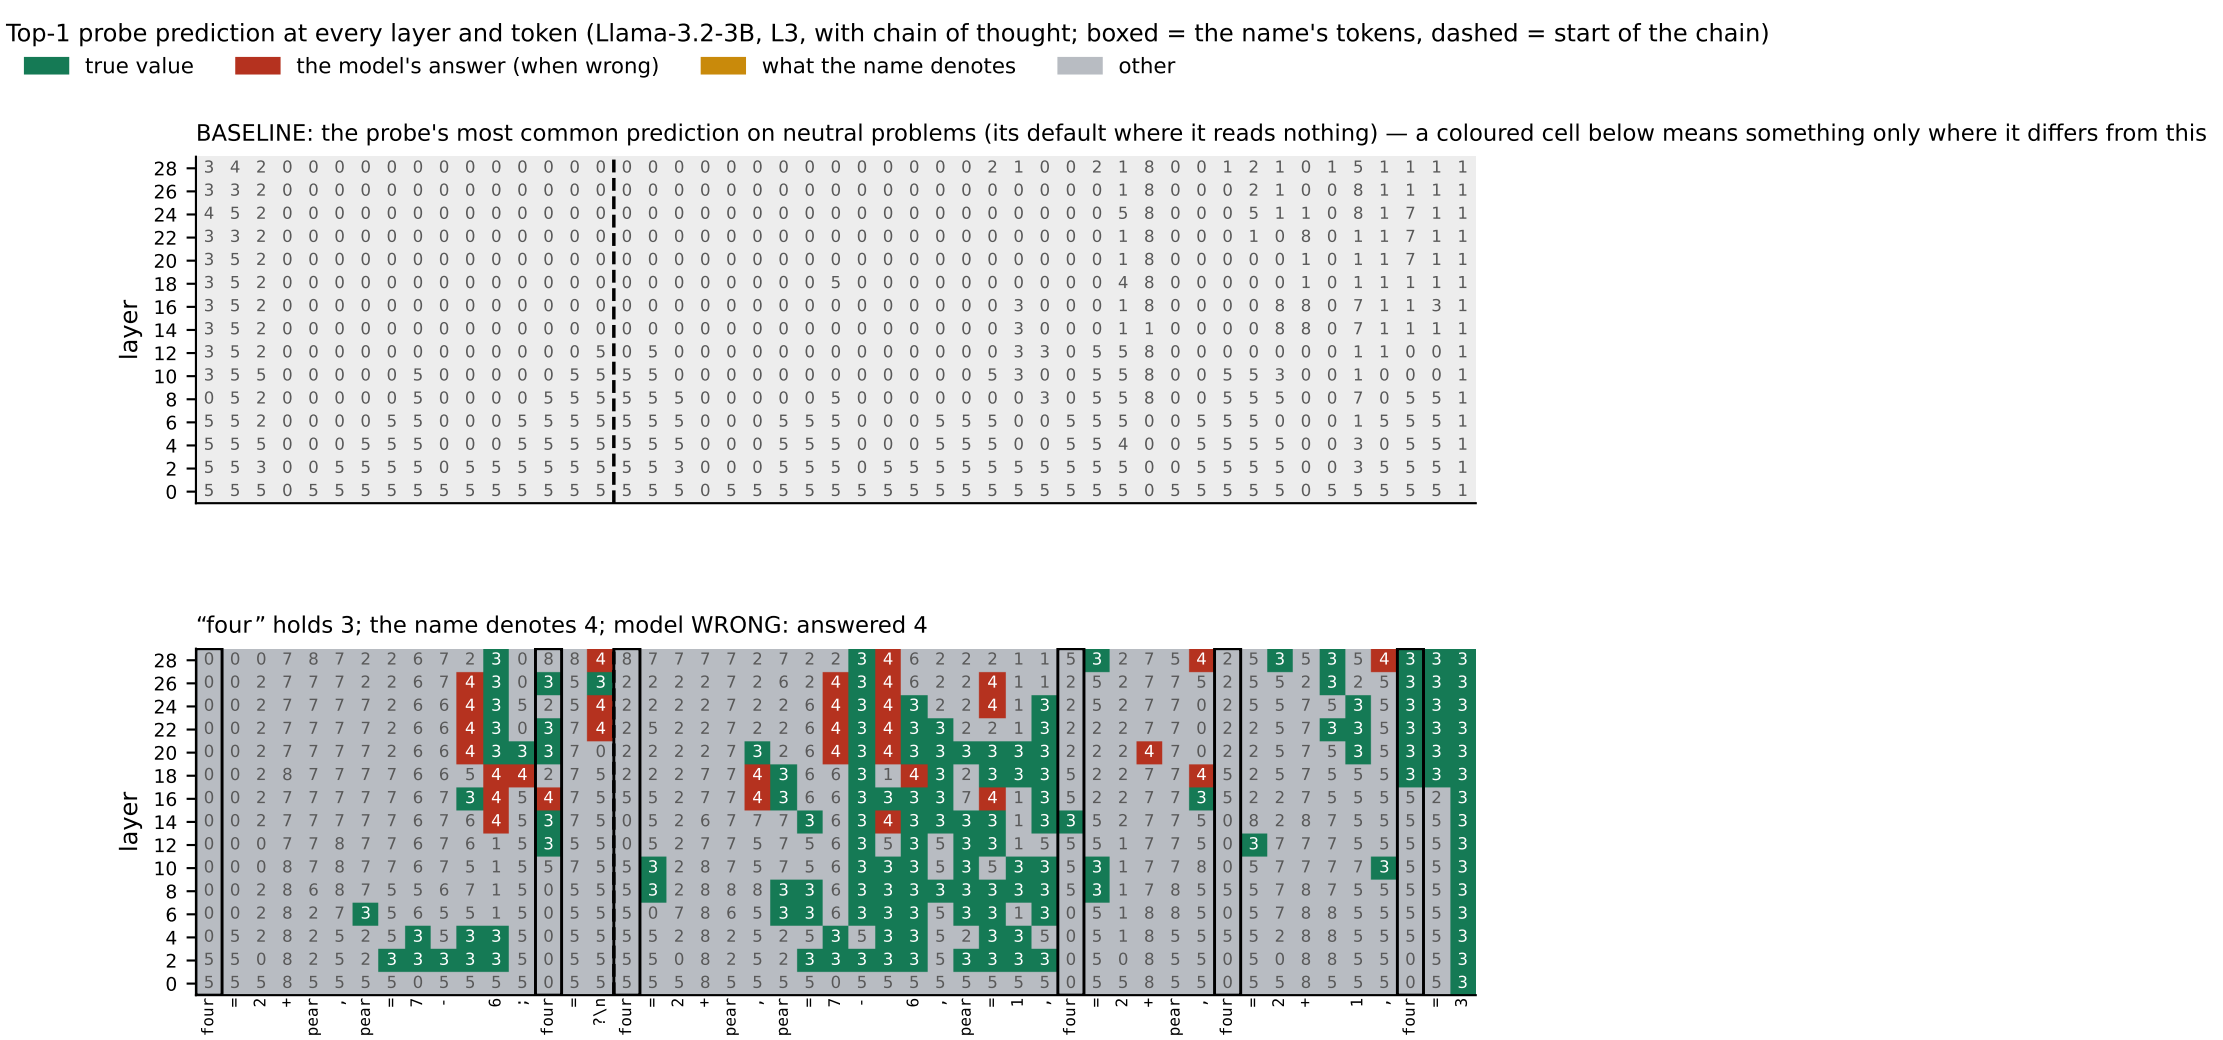

Top-1 probe prediction at every layer and token (Llama-3.2-3B, L3, with chain of thought; boxed = the name's tokens, dashed = start of the chain)
true value
the model's answer (when wrong)
what the name denotes
other
BASELINE: the probe's most common prediction on neutral problems (its default where it reads nothing) — a coloured cell below means something only where it differs from this
layer
“four” holds 3; the name denotes 4; model WRONG: answered 4
layer