


Top-1 probe prediction at every layer and token (Llama-3.2-3B, L3, with chain of thought; boxed = the name's tokens, dashed = start of the chain)

■ true value    ■ the model's answer (when wrong)    ■ what the name denotes    ■ other



“four” holds 2; the name denotes 4; model WRONG: answered 4

