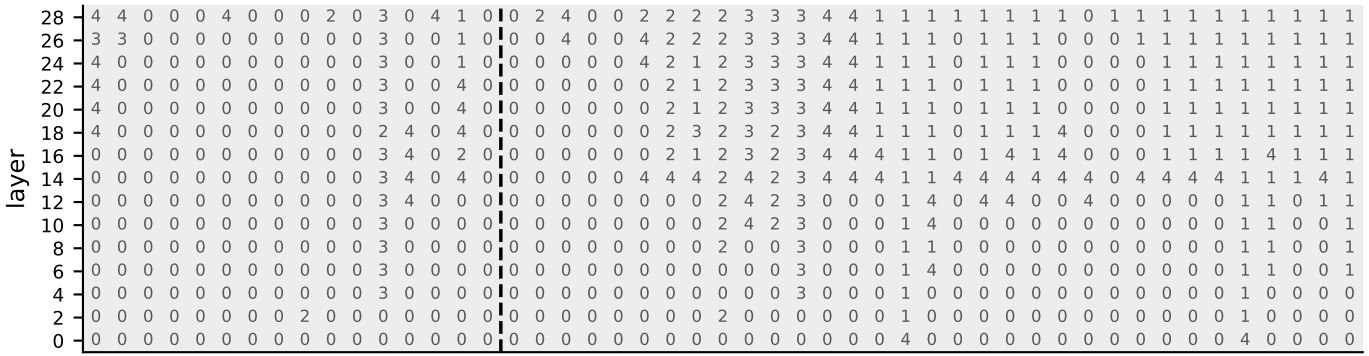
Top-1 probe prediction at every layer and token (Llama-3.2-3B, L3, with chain of thought; boxed = the name's tokens, dashed = start of the chain)

true value | the model's answer (when wrong) | what the name denotes | other

BASELINE: the probe's most common prediction on neutral problems (its default where it reads nothing) — a coloured cell below means something only where it differs from this



“one” holds 3; the name denotes 1; model WRONG: answered 3

