









llama31-8b, level 3, cot: per-token probes for the value of v1 (trained on train_neutral)



control-task accuracy (Hewitt & Liang)



control-task accuracy (Hewitt & Liang)



margin log p(true) - log p(lure)

