



llama31-8b, level 3, direct: per-token probes for the value of v1 (trained on train_neutral)





neutral: value of v1




congruent@v1: value of v1




incongruent@v1: value of v1



control-task accuracy (Hewitt & Liang)



control-task accuracy (Hewitt & Liang)



margin log p(true) - log p(lure)

