



llama31-8b, level 3, direct: per-token probes for the value of v2 (trained on train_neutral)






neutral: value of v2




congruent@v2: value of v2




incongruent@v2: value of v2



control-task accuracy (Hewitt & Liang)



control-task accuracy (Hewitt & Liang)



margin log p(true) - log p(lure)

