



llama32-3b, level 3, direct: per-token probes for the value of v2 (trained on train_letter)





neutral: value of v2



congruent@v2: value of v2




incongruent@v2: value of v2




probe accuracy

control-task accuracy (Hewitt & Liang)



control-task accuracy (Hewitt & Liang)



margin log p(true) - log p(lure)

